

Navnith Bharadwaj

720-725-3458 | bharadwajnavnith5@gmail.com | linkedin.com/in/navnithbharadwaj | github.com/navnithb2001

EDUCATION

University of Colorado Boulder – Aug 2025 – May 2027

Master of Science in Computer Science – GPA 3.74 / 4.00 – Boulder, Colorado

- Relevant Coursework: Big Data, Deep Learning, Neural Networks, NLP, Computer Vision

R V College of Engineering – Aug 2019 – Jul 2023

Bachelor of Engineering in Computer Science and Engineering – Bengaluru, India

- Relevant Coursework: Data Structures, Algorithms, Operating Systems, Software Engineering, DBMS, OOP

TECHNICAL SKILLS

Languages: Python, TypeScript, JavaScript, Java, C++, SQL, Bash

Frontend: React, Tailwind, Vite, WebSockets, Server-Sent Events

Backend and APIs: FastAPI, Spring Boot, REST, GraphQL, gRPC, microservices, PostgreSQL, Redis

Cloud and Infrastructure: Kubernetes, Docker, AWS, Azure, GCP, CI/CD, Git, Linux, monitoring

EXPERIENCE

AI Software Engineer Intern – Aug 2026 – Present

Lyntra – Remote

- Shipped **Gmail and Sheets integrations** converting email into **coursework tasks**, cutting professor setup time by **60%**.

Software Developer – Aug 2023 – Aug 2025

Oracle – Bengaluru, India

- Led a **Backup and Disaster Recovery** automation platform spanning VM backup, database backup and recovery, and VM DR end-to-end, directing 1 engineer, adopted by **3 infra teams**.
- Designed and shipped 35 RBAC-secured **REST APIs** and **GraphQL** schemas standardizing backup and recovery operations, cutting average request-to-execution time from **6 hours to 5 minutes**.
- Built **asynchronous execution pipelines** for long-running backup and recovery jobs, cutting user-facing wait times by **45%** for **high-volume data processing**.
- Instituted CI/CD pipelines, peer **code reviews**, and unit and integration tests at **85% coverage**, cutting production rollback incidents by **35%**.

Backend Developer Intern – Sept 2022 – Nov 2022

Ideoholics – Remote

- Shipped **Java** and Spring Boot REST APIs within a microservices architecture, sustaining **200,000 requests per day** at **99.5% uptime**.

PROJECTS

MeshML: Distributed Platform for ML Training | Python, PyTorch, Kubernetes, gRPC, Docker

- Architected a **distributed platform** to train a shared **PyTorch** model via a parameter server over **6 gRPC services**.
- Stabilized **async data-parallel SGD** across **4 workers** with **staleness-aware** aggregation and server-side momentum.
- Built a **React and Tailwind dashboard** streaming live training metrics over **WebSockets**, holding **sub-65ms p95** API latency.

ML Model A/B Testing Framework | Python, FastAPI, scikit-learn, PostgreSQL

- Engineered a **FastAPI** framework to A/B test **champion vs challenger** models with hot-swappable, **zero-downtime** versions.
- Computed **accuracy, F1, AUC**, p95/p99 latency with Welch's t-test and z-test over **1,000** predictions to flag **regressions**.
- Automated **promote/keep/review** decisions via **95% confidence intervals** and **consistent-hashing** traffic routing.

F1 Telemetry-to-Narrative: AI Agent Pipeline | Python, MLX, OpenAI, Pydantic

- Designed a **6-stage LLM pipeline** turning 50 telemetry frames/session into commentary via local **MLX** and **OpenAI**.
- Ran **A/B comparison** of local vs frontier-LLM outputs and built a **2,000-pair SFT dataset** for **fine-tuning**.
- Enforced reproducibility with **Pydantic-validated** JSONL traces and full ID linkage, streamed over **Server-Sent Events**.